



deeplearning.ai

Introduction to ML strategy

Why ML Strategy?

Motivating example



Ideas:

- Collect more data ←
- Collect more diverse training set
- Train algorithm longer with gradient descent
- Try Adam instead of gradient descent
- Try bigger network
- Try smaller network
- Try dropout
- Add L_2 regularization
- Network architecture
 - Activation functions
 - # hidden units
 - ...




deeplearning.ai


Introduction to ML strategy

Orthogonalization

TV tuning example



Orthogonalization




Car




→ Steering]

→ { Acceleration
Braking }

$$\rightarrow \frac{0.3 \times \text{angle} - 0.8 \text{ speed}}{2 \times \text{angle} + 0.9 \text{ speed}}$$



Chain of assumptions in ML






deeplearning.ai

Setting up
your goal

Single number
evaluation metric

Using a single number evaluation metric



→ Of examples recognized as cat,
what % actually are cats?
→ what % of actual cats
are correctly recognized


	Classifier	Precision	Recall
→	A	95%	90%
→	B	98%	85%

F₁ score = "Average" of P and R.

$$\left(\frac{2}{\frac{1}{P} + \frac{1}{R}} \right) \text{ "Harmonic mean"}$$

Dev set + Single number evaluation metric
real speed up iterating

Another example



Algorithm	US	China	India	Other
A	<u>3%</u>	7%	5%	9%
B	5%	6%	5%	10%
C	2%	3%	4%	5%
D	5%	8%	7%	2%
E	4%	5%	2%	4%
F	7%	11%	8%	12%



deeplearning.ai

Setting up
your goal

Satisficing and
optimizing metrics

Another cat classification example

optimizing ↓ *satisficing* ↓

Classifier	Accuracy	Running time
A	90%	<u>80ms</u>
B	<u>92%</u>	<u>95ms</u>
C	95%	<u>1,500ms</u>

←

$$\text{Cost} = \text{accuracy} - 0.5 \times \text{Running Time}$$

maximize accuracy

subject to Running Time $\leq$ 100 ms.

N metrics : 1 optimizing
N-1 satisficing

Wakewords / Trigger words

Alexa, OK Google,

Hey Siri, nihao baidu
你好 百度

accuracy.

#false positive

maximize accuracy.

s.t. $\leq$ 1 false positive
every 24 hours.



deeplearning.ai

Setting up
your goal

Train/dev/test
distributions

Cat classification dev/test sets

development set, hold out cross validation set

Regions:

- US
- UK
- Other Europe
- South America
- India
- China
- Other Asia
- Australia


Dev

Test

Randomly shuffle into dev/test



dev set
+
metric



True story (details changed)

[Optimizing on dev set on loan approvals for
medium income zip codes




$x \rightarrow y$ (repay loan?)




[Tested on low income zip codes

~ 3 month



Guideline

Choose a dev set and test set to reflect data you expect to get in the future and consider important to do well on.






deeplearning.ai

Setting up
your goal

Size of dev
and test sets

Old way of splitting data



Size of dev set

A B

Set your dev set to be big enough to detect differences in
algorithm/models you're trying out.

100 : small
 ↳ 1%

1,000

10,000

100,000


^A 97% → ^B 97.1%
 0.1%
 ↑

↙ 0.01%
 0.001%

On the advert size

Size of test set

- Set your test set to be big enough to give high confidence in the overall performance of your system.





deeplearning.ai

Setting up
your goal

When to change
dev/test sets and
metrics

Cat dataset examples

Metric + Dev : Prefer A
You/users : Prefer B.

→ Metric: classification error

Algorithm A: 3% error

→ pornographic

✓ Algorithm B: 5% error

$$\left\{ \begin{array}{l} \text{Error: } \frac{1}{\sum_i w^{(i)}} \cdot \frac{1}{m_{\text{dev}}} \sum_{i=1}^{m_{\text{dev}}} w^{(i)} \mathbb{I}\{y_{\text{pred}}^{(i)} \neq y^{(i)}\} \\ \rightarrow w^{(i)} = \begin{cases} 1 & \text{if } x^{(i)} \text{ is non-porn} \\ 10 & \text{if } x^{(i)} \text{ is porn} \end{cases} \end{array} \right.$$

$\mathbb{I}\{y_{\text{pred}}^{(i)} \neq y^{(i)}\}$
predicted value (0/1)

Orthogonalization for cat pictures: anti-porn

- 1. So far we've only discussed how to define a metric to evaluate classifiers. ← Place target ↗
- 2. Worry separately about how to do well on this metric. ↗
↖ Aim (shoot at target)

$$\rightarrow J = \frac{1}{\sum w^{(i)}} \sum_{i=1}^m w^{(i)} \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$



Another example


Algorithm A: 3% error

✓ Algorithm B: 5% error ←

→ Dev/test



→ User images



If doing well on your metric + dev/test set does not correspond to doing well on your application, change your metric and/or dev/test set.




deeplearning.ai

Comparing to human-level performance

Why human-level performance?

Comparing to human-level performance



Why compare to human-level performance

Humans are quite good at a lot of tasks. So long as ML is worse than humans, you can:

- - Get labeled data from humans. (x, y)
- - Gain insight from manual error analysis:
Why did a person get this right?
- - Better analysis of bias/variance.






deeplearning.ai

Comparing to human-
level performance

Avoidable bias

Bias and Variance



Bias and Variance

Cat classification

Human-level $\approx 0\%$

Training set error:

Dev set error:




high variance

high bias

high bias
high variance

low bias
low variance

Cat classification example





deeplearning.ai

Comparing to human-
level performance

Understanding
human-level
performance

Human-level error as a proxy for Bayes error

Medical image classification example:



Suppose:

(a) Typical human 3 % error

→ (b) Typical doctor 1 % error


(c) Experienced doctor 0.7 % error

→ (d) Team of experienced doctors .. 0.5 % error ←


Bayes error $\leq$ 0.5 %

What is “human-level” error?

Error analysis example



Summary of bias/variance with human-level performance





deeplearning.ai

Comparing to human-
level performance

Surpassing human-
level performance

Surpassing human-level performance

Team of humans

0.5%

One human

0.1

~~1.0%~~

Training error


0.6%

Dev error

0.2

0.8%

What is avoidable bias?



Problems where ML significantly surpasses human-level performance

- - Online advertising
- - Product recommendations
- - Logistics (predicting transit time)
- - Loan approvals

Structured data

Not natural perception

Lots of data

- Speech recognition
- Some image recognition
- Medical
 - ECG, Skin cancer, ...



deeplearning.ai

Comparing to human-
level performance

Improving your model
performance

The two fundamental assumptions of supervised learning

1. You can fit the training set pretty well.



~ Avoidable bias

2. The training set performance generalizes pretty well to the dev/test set.



~ Variance

Reducing (avoidable) bias and variance

