



deeplearning.ai

Object Detection


Object
localization

What are localization and detection?

Image classification



Classification with
localization




Detection



"Car"
1 object

multiple
objects


Classification with localization



Defining the target label y

- 1 - pedestrian
- 2 - car ←
- 3 - motorcycle
- 4 - background ←

Need to output b_x, b_y, b_h, b_w , class label (1-4)



$x =$

$$L(\hat{y}, y) = \begin{cases} (\hat{y}_1 - y_1)^2 + (\hat{y}_2 - y_2)^2 + \dots + (\hat{y}_8 - y_8)^2 & \text{if } \underline{y_1 = 1} \\ (\hat{y}_1 - y_1)^2 & \text{if } \underline{y_1 = 0} \end{cases}$$

$y =$

$$\begin{bmatrix} p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \end{bmatrix}$$

is there any object?

(x, y)

$$\begin{bmatrix} 1 \\ b_x \\ b_y \\ b_h \\ b_w \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

$$\begin{bmatrix} 0 \\ ? \\ ? \\ ? \\ ? \\ ? \\ ? \\ ? \end{bmatrix}$$

p_c
← "don't care"



deeplearning.ai




Object Detection

Landmark
detection

Landmark detection



b_x, b_y, b_h, b_w



$l_{1x}, l_{1y},$
 $l_{2x}, l_{2y},$
 $l_{3x}, l_{3y},$
 $l_{4x}, l_{4y},$
 $\vdots$
 l_{64x}, l_{64y}

X, y

$l_{1x}, l_{1y},$
 $\vdots$
 l_{32x}, l_{32y}



deeplearning.ai

Object Detection

Object
detection

Car detection example

Training set:

x

y



1



1



1



0




0




$\rightarrow$ ConvNet $\rightarrow y$

Sliding windows detection


→ ConvNet → 0



→ ConvNet



Computation cost






deeplearning.ai


Object Detection

Convolutional
implementation of
sliding windows

Turning FC layer into convolutional layers




Convolution implementation of sliding windows



[Sermanet et al., 2014, OverFeat: Integrated recognition, localization and detection using convolutional networks]

Andrew Ng

Convolution implementation of sliding windows






deeplearning.ai


Object Detection

**Bounding box
predictions**

Output accurate bounding boxes






YOLO algorithm



Labels for training
For each grid cell:

Target output:


$3 \times 3 \times 8$



[Redmon et al., 2015, You Only Look Once: Unified real-time object detection]

Andrew Ng

Specify the bounding boxes






deeplearning.ai

Object Detection

Intersection
over union

Evaluating object localization



Intersection over Union (IoU)

$$= \frac{\text{size of } \text{[yellow hatched box]}}{\text{size of } \text{[green hatched box]}}$$

“Correct” if $\text{IoU} \geq 0.5$ ←

0.6 ←

More generally, IoU is a measure of the overlap between two bounding boxes.



deeplearning.ai


Object Detection

Non-max
suppression

Non-max suppression example




Non-max suppression example




19x19

Non-max suppression example




p_c

Non-max suppression algorithm



19×19

Each output prediction is:



Discard all boxes with $p_c \leq 0.6$

→ While there are any remaining boxes:

- Pick the box with the largest p_c .
Output that as a prediction.
- Discard any remaining box with $\text{IoU} \geq 0.5$ with the box output in the previous step

Andrew Ng




deeplearning.ai

Object Detection

Anchor boxes


Overlapping objects:




$$y = \begin{bmatrix} p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \end{bmatrix}$$

Handwritten annotations: A green arrow points from p_c to the center of the woman's bounding box. A blue arrow points from b_x to the left edge of the woman's bounding box. A blue bracket groups c_1, c_2, c_3 .

Anchor box 1:



Anchor box 2:



$y =$


Anchor box algorithm

Previously:

Each object in training image is assigned to grid cell that contains that object's midpoint.

Output y :

$$\underline{3 \times 3 \times 8}$$



With two anchor boxes:


Each object in training image is assigned to grid cell that contains object's midpoint and anchor box for the grid cell with highest IoU.

(grid cell, anchor box)



Output y :

$$\begin{array}{l} 3 \times 3 \times \underline{16} \\ 3 \times 3 \times \underline{2 \times 8} \end{array}$$

Anchor box example



Anchor box 1: Anchor box 2:



$y =$

$$\begin{bmatrix} p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \\ p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \end{bmatrix}$$

Handwritten annotations for the first vector y :

- Orange text: $1, b_x, b_y, b_h, b_w$
- Green text: $1, 0, 0, 1, 0$

Handwritten annotations for the second vector:

- Blue text: c_1, c_2, c_3
- Orange text: $?, ?, ?$
- Green text: $1, b_x, b_y, b_h, b_w, 0, 1, 0$
- Blue text: "only?"
- Blue text: "anchor box 1" (pointing to the first vector)
- Blue text: "anchor box 2" (pointing to the second vector)




deeplearning.ai

Object Detection


Putting it together:
YOLO algorithm

Training

- 1 - pedestrian
- 2 - car ←
- 3 - motorcycle



$y =$



$$\begin{bmatrix} p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \\ p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \end{bmatrix}$$

$$\begin{bmatrix} 0 \\ ? \\ ? \\ ? \\ ? \\ ? \\ ? \\ ? \\ 0 \\ ? \\ ? \\ ? \\ ? \\ ? \\ ? \\ ? \end{bmatrix}$$


$$\begin{bmatrix} 0 \\ ? \\ ? \\ ? \\ ? \\ ? \\ ? \\ ? \\ 1 \\ b_x \\ b_y \\ b_h \\ b_w \\ 0 \\ 1 \\ 0 \end{bmatrix}$$

y is $3 \times 3 \times 2 \times 8$

$19 \times 19 \times 16$
 $19 \times 19 \times 40$

#anchors

$5 + \#classes$




ConvNet




Andrew Ng

Making predictions



Outputting the non-max suppressed outputs



- For each grid cell, get 2 predicted bounding boxes.
- Get rid of low probability predictions.
- For each class (pedestrian, car, motorcycle) use non-max suppression to generate final predictions.





deeplearning.ai


Object Detection

Region proposals (Optional)

Region proposal: R-CNN



↑



Segmentation algorithm
~2,000

[Girshik et. al, 2013, Rich feature hierarchies for accurate object detection and semantic segmentation] Andrew Ng

Faster algorithms

→ R-CNN: Propose regions. Classify proposed regions one at a time. Output label + bounding box. ←

Fast R-CNN: Propose regions. Use convolution implementation of sliding windows to classify all the proposed regions. ←

Faster R-CNN: Use convolutional network to propose regions.

[Girshik et. al, 2013. Rich feature hierarchies for accurate object detection and semantic segmentation]

[Girshik, 2015. Fast R-CNN]

[Ren et. al, 2016. Faster R-CNN: Towards real-time object detection with region proposal networks]

Andrew Ng



deeplearning.ai


Convolutional Neural Networks

Semantic segmentation with U-Net


Object Detection vs. Semantic Segmentation



Input image




Object Detection




Semantic Segmentation

Motivation for U-Net



Chest X-Ray




Brain MRI

[Novikov et al., 2017, Fully Convolutional Architectures for Multi-Class Segmentation in Chest Radiographs]

[Dong et al., 2017, Automatic Brain Tumor Detection and Segmentation Using U-Net Based Fully Convolutional Networks]


Andrew Ng

Per-pixel class labels




1. Car
0. Not Car

Per-pixel class labels




1. Car
2. Building
3. Road




Segmentation Map

Deep Learning for Semantic Segmentation




Transpose Convolution


Normal Convolution




Transpose Convolution




Transpose Convolution



Deep Learning for Semantic Segmentation




U-Net



[Ronneberger et al., 2015, U-Net: Convolutional Networks for Biomedical Image Segmentation]

Andrew Ng

U-Net



[Ronneberger et al., 2015, U-Net: Convolutional Networks for Biomedical Image Segmentation]

Andrew Ng